UMANG PAWAR

AI/GenAI Engineer (Fresher) | LLMs, RAG & Agentic Systems
Nagpur, India | +91 8055070449 | umangpawar629@gmail.com
linkedin.com/in/umang-pawar-208157314 | github.com/Umang07-cell | umangpawar.vercel.app

SUMMARY

Fresher AI Engineer who built and deployed RAG systems, guardrailed retrieval, and agentic LLM workflows across 4 AI projects and 2 live deployments. Built an access-controlled HR RAG chatbot for 100+ employees at LanceSoft (45/50 on a hand-checked test set). Skilled in Python, FastAPI, LangChain, and 3 vector stores (Qdrant, FAISS, pgvector).

SKILLS

AI / GenAI: Retrieval-Augmented Generation (RAG), LangChain, Agentic Workflows, LoRA Fine-tuning, Prompt Engineering, Reranking, Guardrails, LLM Evaluation, Function/Tool Calling, Vision-Language Models
LLMs / Vector DBs: Groq, OpenAI API, Llama 3, Qdrant, FAISS, pgvector
Languages / Backend: Python, SQL, FastAPI, Flask, REST APIs, SQLAlchemy, ABAC (Attribute-Based Access Control)
Data / Infra: PostgreSQL, Redis, Pandas, Power BI, Excel, Git, GitHub Actions (CI/CD), Docker, Railway, Render

PROFESSIONAL EXPERIENCE

AI Engineer Intern | LanceSoft **Jul 2026 - Sep 2026**

- Built **HR-PolicyIQ**, a RAG chatbot (LangChain, Qdrant, Groq Llama 3, FastAPI, PostgreSQL, Redis) over **30+ HR policy documents (~4,000 chunks)**, deployed to **100+ employees**.
- Evaluated on a **50-question** test set, checking every answer against source documents: **45/50 correct** with relevant citations; fixed the 5 partially accurate answers by revising the system prompt.
- Implemented ABAC retrieval for **3 roles** (manager, HR, employee) across Hyderabad and Bangalore; ran **20+ edge-case** unauthorized-access tests, identified content leaks, and fixed them.
- Added a confidence-scored reranker and source-cited answers as grounding guardrails against hallucination, plus request logging to monitor query and API-key failures.
- Benchmarked local Llama 3 (PyTorch) against Groq-hosted Llama 3 and migrated to Groq for production reliability; cut average response latency from 5-6s to 1-1.5s (**~4x faster**) via embedding caching and pipeline optimization.
- Built a lead-generation engine matching **3,000+ candidate records** to job descriptions and resumes using pgvector and contact retrieval, surfacing qualified leads for managers and HR.

Data Analyst Intern | CarrierDomain **Mar 2026 - Jun 2026**

- Built automated Python/Pandas/SQL ETL pipelines processing **637,494+ client records** across multiple source tables into ML-ready datasets.
- Built Power BI and Excel dashboards that refresh directly from the pipeline, giving stakeholders weekly KPI and segmentation reporting.

Software Developer Intern | Syncsas Technologies **Nov 2025 - Mar 2026**

- Shipped **3 client-facing** full-stack features (Python, Flask, PostgreSQL) across Agile sprints, including an AI chatbot (LangChain + OpenAI API), with testing.
- Set up GitHub Actions CI/CD pipelines, cutting manual deployment steps by **60%**.

PROJECTS

Atlas-AI: AI Career Platform **2026**

Python, *FastAPI*, *Groq*, *LangChain*, *JSearch* | Live Demo | GitHub

- Built **4 integrated AI systems** in one deployed product: a JSearch job scraper (8-10 jobs per run), an ATS resume scorer, a voice mock-interview engine with LLM follow-ups, and a RAG career chatbot.
- Built an ATS analyzer returning a match score, missing keywords and skills, and improvement suggestions against a job description.
- Fine-tuned Llama 3 with LoRA on a public Kaggle dataset, served for interview follow-ups; integrated a Groq-hosted vision-language model.
- Cut response latency from 4-5s to ~2s (**~2x faster**).

FinSight: AI Financial Intelligence Platform **2026**

Groq, *LangChain*, *FAISS*, *Railway* | Live Demo

- Built a RAG chatbot answering questions on financial reports with cited sources in under 2s.
- Tested on sample financial reports with **20+ edge-case questions**, checking answers against the source; improved answer quality by fixing chunking and tuning the retrieval pipeline.
- Deployed on Railway; resolved Pydantic version conflicts and CORS issues.

Closira: AI Customer Support Workflow **2026**

Python, *Groq API*, *Llama 3.3* | GitHub

- Built a **4-stage** LLM agentic pipeline for SOP-grounded FAQ answering, lead qualification, escalation, and session summarization, with a two-layer escalation check (keyword pre-check plus model-generated tags).

EDUCATION

B.Tech, Computer Science & Engineering | G.H. Raisoni University, Nagpur **2026**

CGPA: 7.50/10

CERTIFICATIONS & PUBLICATION

Google AI Professional Certificate (2026) | Research Publication, IJRASET (2025)